

# Assignments for week 3

Anders Degn Lapiki, Jacob Kjærulff Furberg, Jonas Ishøj Nielsen

03/09/2020

## Contents

<b>1</b>	<b>4.1 Microbenchmarks</b>	<b>2</b>
1.1	4.1 - red . . . . .	2
<b>2</b>	<b>4.2 Standard Deviation</b>	<b>4</b>
2.1	4.2 - green . . . . .	4
2.2	4.2 - yellow . . . . .	5
2.3	4.2 - red . . . . .	5
<b>3</b>	<b>4.3 TestTimer</b>	<b>5</b>
3.1	4.3.1 - green . . . . .	5
3.2	4.3.2 - yellow . . . . .	5
<b>4</b>	<b>4.4 TestCountPrimes</b>	<b>6</b>
4.1	4.4.2 - green . . . . .	6
4.2	4.4.3 - yellow . . . . .	6
4.3	4.4.4 - red . . . . .	6
<b>5</b>	<b>4.5 - Synchronised test</b>	<b>6</b>
5.1	4.5 - green . . . . .	6
<b>6</b>	<b>4.6 Volatile vs. Non-volatile</b>	<b>7</b>
6.1	4.6.1 - green . . . . .	7
6.2	4.6.2 - red . . . . .	7

## 1 4.1 Microbenchmarks

### 1.1 4.1 - red

System

```
# OS: Windows 10; 10.0; amd64\\
# JVM: Oracle Corporation; 14.0.2\\
# CPU: Intel64 Family 6 Model 94 Stepping 3, GenuineIntel; 8 "cores"\\
# Date: 2020-09-17T10:27:20+0200\\
```

Mark1

0,009 s      0,5ns

Very low, as the multiply is treated as dead code by the JVM.

Mark2

23,6 ns

Makes sense, now that the code is used.

Mark3

23,7 ns  
23,5 ns  
23,5 ns  
25,0 ns  
24,3 ns  
23,5 ns  
23,5 ns  
23,5 ns  
23,5 ns  
23,5 ns

Makes sense, now that the code is used and test is run x times to show more reasonable results. The individual result are also close to Mark2 as expected.

Mark4

23,7 ns +/- 0,626

The mean and the standard deviation seems to fit the result when run 10 times.

Mark5

425,0 ns +/-	817,60	2
135,0 ns +/-	67,91	4
115,0 ns +/-	38,55	8
103,1 ns +/-	20,04	16
110,3 ns +/-	43,83	32
91,9 ns +/-	8,46	64
48,8 ns +/-	52,13	128
32,6 ns +/-	2,97	256
31,3 ns +/-	1,35	512
30,9 ns +/-	0,52	1024
30,7 ns +/-	0,26	2048
30,7 ns +/-	0,18	4096
33,1 ns +/-	4,33	8192
30,5 ns +/-	0,36	16384
33,6 ns +/-	20,52	32768
29,8 ns +/-	2,65	65536
26,2 ns +/-	0,51	131072
23,5 ns +/-	0,55	262144
23,5 ns +/-	0,15	524288
23,5 ns +/-	0,15	1048576
23,5 ns +/-	0,13	2097152
23,5 ns +/-	0,18	4194304
23,5 ns +/-	0,06	8388608
23,5 ns +/-	0,11	16777216

Makes sense as the mean converges towards correct answer, standard deviation becomes smaller and garbage collection cause outliers underway.

#### Mark6

multiply	845,0 ns	1882,00	2
multiply	627,5 ns	1071,89	4
multiply	268,8 ns	75,29	8
multiply	268,8 ns	190,92	16
multiply	88,8 ns	79,74	32
multiply	63,8 ns	9,10	64
multiply	113,1 ns	102,51	128
multiply	97,3 ns	49,50	256
multiply	67,1 ns	30,59	512
multiply	34,6 ns	0,76	1024
multiply	36,3 ns	4,14	2048
multiply	35,3 ns	2,23	4096
multiply	32,0 ns	3,98	8192
multiply	29,9 ns	0,31	16384
multiply	29,2 ns	2,28	32768
multiply	30,8 ns	1,85	65536
multiply	25,3 ns	1,50	131072

multiply	24,9 ns	0,47	262144
multiply	23,5 ns	0,14	524288
multiply	23,5 ns	0,18	1048576
multiply	23,5 ns	0,12	2097152
multiply	23,5 ns	0,08	4194304
multiply	23,5 ns	0,06	8388608
multiply	23,5 ns	0,08	16777216

It fits that the Mark6 is similar to Mark5, as Mark6 is just Mark5 but more generalized.

Mark7

23,5 ns      0,26    16777216\\

Not much to say, as the only difference is, that Mark7 only prints the final measurement.

mark7math

pow	19,3 ns	0,03	16777216
exp	21,2 ns	0,05	16777216
log	11,3 ns	0,05	33554432
sin	14,3 ns	0,07	33554432
cos	14,1 ns	0,02	33554432
tan	19,5 ns	0,03	16777216
asin	213,8 ns	0,38	2097152
acos	202,8 ns	0,38	2097152
atan	43,1 ns	0,12	8388608

It fits that count differ as they don't all take the same amount of time. The standard deviation should also be larger for the slower functions which fit.

The differences in our results and those found in Microbenchmarks, can be explained by garbage collection, system, java version, and so on.

## 2 4.2 Standard Deviation

### 2.1 4.2 - green

The mean is the average of the values, and in Gaussian distribution it is represented by  $\mu$ .

The standard deviation sdev is the  $\sigma$  and is how much an average differs from the mean.

The variance would be calculated as:  $\sigma^2$ .

The code calculates the standard deviation. The code looks different from what is shown in the slide, since the formula has been somewhat rewritten.

The completed code can be seen in `MeanVar.java`.

And the output when running on the array: `[30.7, 30.3, 30.1, 30.7, 50.2, 30.4, 30.9, 30.3, 30.5, 30.8]` 32.0 ns +/- 8.576.

And the output when running on the array: `[30.7, 100.2, 30.1, 30.7, 20.2, 30.4, 2, 30.3, 30.5, 5.4]` 30.0 ns +/- 27.974.

## 2.2 4.2 - yellow

Without any 25, the 25 is not an outlier as 25 is within  $|32.0 - 8.576, 30.0 + 8.576|$

The output when running on the array: `[30.7, 100.2, 30.1, 30.7, 20.2, 30.4, 2, 30.3, 30.5, 5.4, 25]` 30.0 ns +/- 26.015.

In this example 25 is not an outlier as 25 is within  $|30.0 - 26.015, 30.0 + 26.015|$ .

## 2.3 4.2 - red

The general purpose java program for calculating mean and standard deviation of a large data file can be found in `MeanVar.java`.

## 3 4.3 TestTimer

### 3.1 4.3.1 - green

The large variation in the time measurement is due to other things than the actual code. The JVM setting up and the likes. Therefore at some point, with enough test cases, the hidden work by the JVM will become insignificant.

### 3.2 4.3.2 - yellow

Using mark7 the result is:

hashCode()	2,5 ns	0,00	134217728
Point creation	65,0 ns	0,98	4194304
Thread's work	5066,3 ns	18,75	65536
Thread create	1094,3 ns	14,96	262144
Thread create start	95605,9 ns	218,66	4096
Thread create start join	222040,8 ns	373,98	2048
ai value = 1433540000			
Uncontended lock	4,8 ns	0,01	67108864

All the results makes sense, as creating a thread takes longer time than creating an object. It also makes sense that starting a thread takes longer time, and joining adds extra time.

## 4 4.4 TestCountPrimes

The result can be seen in 4.4.1.txt and was done by running with 1, 3, 5, ..., 31 threads.

### 4.1 4.4.2 - green

As expected the standard deviation decreases with count and the change between the mean results diminish as count increases. When this isn't true it is possibly due to garbage collection.

When running on a 8 core computer, using 13-15 threads seems optimal.

### 4.2 4.4.3 - yellow

The result can be seen in 4.4.3.txt.

The performance of AtomicLong is better than LongCounter and the optimal thread count has also been increased from 13-15 to 23.

If a general use adequate built-in classes and methods exist, they should be used.

### 4.3 4.4.4 - red

The result can be seen in 4.4.5.txt.

This changed doesn't make the result faster and infact only makes it slower. The result is more similar to that found in 4.4.1.

## 5 4.5 - Synchronised test

### 5.1 4.5 - green

We made a minor change to the code. The code given in the exercises was:

```
Mark6("Uncontended lock",i -> {synchronized (obj) {return i;}});
```

This would not work, so we changed it to:

```
Mark6("Uncontended lock",i -> {synchronized (new int[0]) {return i;}});
```

The results from this code where:

Uncontended lock	1585.0 ns	3485.13	2
Uncontended lock	487.5 ns	330.25	4
Uncontended lock	452.5 ns	239.34	8
Uncontended lock	171.3 ns	86.01	16
Uncontended lock	243.1 ns	224.93	32
Uncontended lock	63.9 ns	18.55	64
Uncontended lock	58.9 ns	12.66	128

Uncontended lock	66.1 ns	14.36	256
Uncontended lock	56.7 ns	7.41	512
Uncontended lock	40.2 ns	3.81	1024
Uncontended lock	39.5 ns	3.98	2048
Uncontended lock	39.1 ns	1.35	4096
Uncontended lock	19.1 ns	14.31	8192
Uncontended lock	11.3 ns	2.70	16384
Uncontended lock	10.4 ns	0.57	32768
Uncontended lock	11.5 ns	0.84	65536
Uncontended lock	4.7 ns	3.21	131072
Uncontended lock	19.5 ns	19.28	262144
Uncontended lock	5.2 ns	2.00	524288
Uncontended lock	5.7 ns	2.21	1048576
Uncontended lock	7.2 ns	1.94	2097152
Uncontended lock	4.3 ns	1.94	4194304
Uncontended lock	4.8 ns	1.12	8388608
Uncontended lock	4.3 ns	1.55	16777216
Uncontended lock	5.0 ns	1.45	33554432
Uncontended lock	4.9 ns	0.91	67108864

These results make sense. There are some outliers due to garbage collection, but otherwise the average execution time goes down with more cases. The number stabilises around 4.5 ns at the end.

## 6 4.6 Volatile vs. Non-volatile

### 6.1 4.6.1 - green

The result can be seen in 4.6.1.txt and was done by adding 2 variables at the top of the class, and calling them in Mark6 in main.

```
static int x1 = 0;
static volatile int x2 = 0;

Mark6("NON-VOLATILE int",i -> {return x1++;});
Mark6("VOLATILE int",i -> {return x2++;});
```

The result indicates that updating a non-volatile int is around 8 times faster than a volatile one.

### 6.2 4.6.2 - red

```
***VOLATILE EXPERIMENTS***
N=10000000
Volatile time: 79682100
Volatile time per iteration: 7.9682097
```

```
N=100000000
  Volatile time: 779600300
  Volatile time per iteration: 7.7960033

N=214748364
  Volatile time: 300
  Volatile time per iteration: 1.3969839E-6

***NON-VOLATILE EXPERIMENTS***
N=100000000
  Non volatile time: 9087900
  Non Volatile time per iteration: 0.90879

N=100000000
  Non volatile time: 12728600
  Non Volatile time per iteration: 0.127286

N=214748364
  Non volatile time: 200
  Non Volatile time per iteration: 9.313226E-7
```

Excluding the last result in both, the times for volatile integers seems similar but smaller than the result of Mark6. Mark6 may take longer time as the result is also added to the dummy variable which mark6 uses. This would also explain why the non-volatile integer test never gets below 1 ns.

The reason why the last result is strange and wrong, we believe is because of the JVM wrongfully trying to optimize the code.