

# Assignments for week 7

Anders Degn Lapiki, Jacob Kjærulff Furberg, Jonas Ishøj Nielsen

01/10/2020

## Contents

<b>1</b>	<b>7.1</b>	<b>2</b>
1.1	7.1.1 - green . . . . .	2
<b>2</b>	<b>7.2</b>	<b>3</b>
2.1	7.2.1 - Green - Get . . . . .	3
2.2	7.2.2 - Green - Size . . . . .	3
2.3	7.2.3 - Green - putIfAbsent . . . . .	3
2.4	7.2.4 - Yellow - putIfAbsent call reallocateBuckets . . . . .	3
2.5	7.2.5 - Yellow - remove . . . . .	3
2.6	7.2.6 - Yellow - forEach . . . . .	3
2.7	7.2.7 - Yellow . . . . .	3
2.8	7.2.8 - Yellow . . . . .	3
2.9	7.2.9 - Red . . . . .	3
2.10	7.2.10 - Red . . . . .	4
2.11	7.2.11 - Red . . . . .	4

## 1 7.1

### 1.1 7.1.1 - green

We chose to make the `ArrayList` concurrent by implementing the `ReentrantReadWriteLock` on both the items and the size of the arraylist. Since the `ReentrantReadWriteLock` uses the methods `lock` and `unlock`, we didn't have to use the `synchronized` keyword for locking.

We considered using a `volatile long/int` for the size, but we chose to "play it safe" and avoid using it.

We considered using striping to split locks all over the item array, similar to how the `ConcurrentHashMap` split its hashmap into smaller parts, where a single lock is responsible for each part. To test it we made an additional version using striping.

We also considered using `LongAdder` objects, instead of `Longs`. This would allow several threads to access each long at the same time. We didn't do it since we return values and not references so an example where a user writes: `long x = lst.get(i);` and then `x=2;` will not affect `lst.get(i)`.

Using `mark6` the scaling performance can be seen in `out71.txt`. The last results are:

msg	mean	sdev	count
Better add:	144493.3 ns	2327.52	4096
Stripping add:	201844.0 ns	15726.56	2048
Synchronized add:	55787.2 ns	4630.80	8192
Better get:	32.6 ns	1.48	8388608
Stripping get:	53.1 ns	1.98	8388608
Synchronized get:	7.5 ns	0.43	33554432
Better set:	30.9 ns	1.42	8388608
Stripping set:	52.9 ns	2.61	8388608
Synchronized set:	7.6 ns	0.21	67108864

In general it would be easier to make the class efficiently synchronized, if we had more information about it. Mainly it would be good to know how it will be used. Does it have many writes? Or is it mainly reads? Also it would be good to know a performance requirement for the class. Such that we knew when to stop modifying the class.

## **2 7.2**

### **2.1 7.2.1 - Green - Get**

Done

### **2.2 7.2.2 - Green - Size**

If the strip `s` is not locked, visibility is not guaranteed. In `StripedWriteMap` `size` is an `AtomicIntegerArray` meaning visibility is guaranteed and thus locking is not required.

### **2.3 7.2.3 - Green - putIfAbsent**

Done

### **2.4 7.2.4 - Yellow - putIfAbsent call reallocateBuckets**

Done

### **2.5 7.2.5 - Yellow - remove**

Done

### **2.6 7.2.6 - Yellow - forEach**

We implemented version (2) as it only locks on 1 out of 16 stripes at a time. Had there been a requirement that they must be iterated over in a specific order then we could not have used this version.

### **2.7 7.2.7 - Yellow**

Using small number reduces memory usage and operations like `size` where all locks are used become constant time instead of linear time.

### **2.8 7.2.8 - Yellow**

Increasing lock count to 32 increases lock granularity. This can reduce lock contention thus giving it better concurrent performance.

### **2.9 7.2.9 - Red**

Result can be seen in `out72.txt`. All four classes give the same result. This is summarized to:

```
class SynchronizedMap, class StripedMap, class StripedWriteMap, class WrapConcurrentHashMap
17 maps to B
117 maps to C
34 maps to F
217 maps to E
17 maps to B
17 maps to B
217 maps to E
34 maps to F
217 maps to E
17 maps to B
34 maps to F
```

## 2.10 7.2.10 - Red

It makes sense that the SynchronizedMap is the slowest and that StripedWriteMap is faster than StripedMap.  
We expected WrapConcHashMap to be faster as it uses java's ConcurrentHashMap implementation, but it was the second slowest

Todo: ask what the ??? column is because it isn't mentioned in code, appears to be a secret print statement called from somewhere.

```
# OS: Mac OS X; 10.15.4; x86_64
# JVM: Oracle Corporation; 14.0.2
# CPU: null; 8 "cores"
# Date: 2020-10-17T11:37:00+0200
```

	threadCount	mean	sdev	count
SynchronizedMap	16	649864,7 us	15359,39	2
99992.0	???			
StripedMap	16	152497,4 us	57897,03	2
99992.0	???			
StripedWriteMap	16	93043,8 us	20696,53	4
99992.0667	???			
WrapConcHashMap	16	250045,7 us	47728,03	2
99992.0	???			

## 2.11 7.2.11 - Red

index of  $k1 = 5 \% 3 = 2$ , index of  $k2 = 8 \% 3 = 2$ .  
index  $0 = 0 \% 2 = \text{stripe0}$ , index  $1 = 1 \% 2 = \text{stripe1}$ , index  $2 = 2 \% 2 = \text{stripe0}$ .

It hurts performance that they don't all guard same amount of buckets, but we don't see any reason why thread safety should suffer.